Application of Deep Learning to Text and Image Data

Module 1, Lab 4: Introducing CNNs

In the previous labs, you used neural networks to predict the target field of a given dataset. You used a feed-forward neural network for a multiclass classification task using images as inputs.

Now you will use a convolutional neural network (CNN) that is specialized to extract useful information from images. You will train and evaluate this network on a dataset of handwritten digits, and you will try to predict a number that is represented in an image.

You will learn how to do the following:

- Build a CNN.
- Train a CNN.
- Test the performance of a CNN.

You will be presented with two kinds of exercises throughout the notebook: activities and challenges.

No coding is needed for an activity. You try to understand a concept, answer questions, or run a code cell.

Challenges are where you can practice your coding skills.

Index

- MNIST dataset
- Creating a CNN
- Training the network
- Testing the network

MNIST dataset

The MNIST dataset is a large collection of handwritten digits. Each example contains a pixel map showing how a person wrote a digit. The images have been size-normalized and centered with fixed dimensions. The labels correspond to the digit in the image, ranging from 0 to 9. This is a multiclass classification task with 10 output classes.

```
%%capture
# Install libraries
!pip install -U -q -r requirements.txt
# Import the library dependencies
import boto3
import os
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import torch
from torch import nn
import torchvision
from torchvision import transforms
from torchvision.datasets import ImageFolder
from torch.optim import SGD
Matplotlib is building the font cache; this may take a moment.
Matplotlib is building the font cache; this may take a moment.
# Load the train data (it's included in the torchvision library)
train data = torchvision.datasets.MNIST(
    root="data", train=True, transform=transforms.ToTensor(),
download=True
# Load the test data (it's included in the torchvision library)
test data = torchvision.datasets.MNIST(
    root="data", train=False, transform=transforms.ToTensor(),
download=True
# Print the dimensions of the datasets
print(
    "Training data shape: {}. \nTest data shape: {}".format(
        list(train data.data.shape), list(test data.data.shape)
)
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-
ubyte.gz
Failed to download (trying next):
HTTP Error 404: Not Found
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/train-
```

```
images-idx3-ubyte.gz
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/train-
images-idx3-ubyte.gz to data/MNIST/raw/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-
ubyte.gz
Failed to download (trying next):
HTTP Error 404: Not Found
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/train-
images-idx3-ubyte.gz
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/train-
images-idx3-ubyte.gz to data/MNIST/raw/train-images-idx3-ubyte.gz
100% | 9912422/9912422 [00:00<00:00, 123834903.25it/s]
Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw
Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-
ubyte.gz
Failed to download (trying next):
HTTP Error 404: Not Found
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/train-
labels-idx1-ubvte.gz
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/train-
labels-idx1-ubyte.gz to data/MNIST/raw/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-
ubvte.qz
Failed to download (trying next):
HTTP Error 404: Not Found
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/train-
labels-idx1-ubvte.gz
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/train-
labels-idx1-ubyte.gz to data/MNIST/raw/train-labels-idx1-ubyte.gz
100% | 28881/28881 [00:00<00:00, 15966217,72it/s]
Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Failed to download (trying next):
HTTP Error 404: Not Found
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/t10k-images-
idx3-ubyte.gz
Downloading https://ossci-datasets.s3.amazonaws.com/mnist/t10k-images-
idx3-ubyte.gz to data/MNIST/raw/t10k-images-idx3-ubyte.gz
```

0% | 0/1648877 [00:00<?, ?it/s]

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz Failed to download (trying next):

HTTP Error 404: Not Found

Downloading https://ossci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz

Downloading https://ossci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz to data/MNIST/raw/t10k-images-idx3-ubyte.gz

100% | 1648877/1648877 [00:00<00:00, 232826939.02it/s]

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz Failed to download (trying next):

HTTP Error 404: Not Found

Downloading https://ossci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz Failed to download (trying next):

HTTP Error 404: Not Found

Downloading https://ossci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz

Downloading https://ossci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz to data/MNIST/raw/t10k-labels-idx1-ubyte.gz

Downloading https://ossci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz to data/MNIST/raw/t10k-labels-idx1-ubyte.gz

100%| 4542/4542 [00:00<00:00, 2187955.53it/s]

Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

Training data shape: [60000, 28, 28].

Test data shape: [10000, 28, 28]

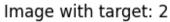
Training data shape: [60000, 28, 28].

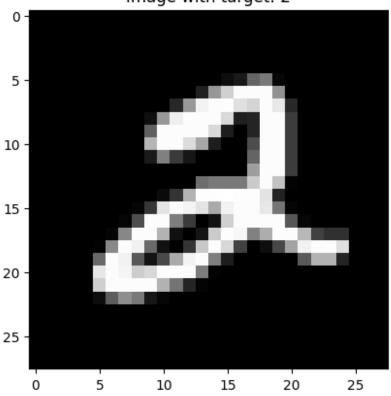
Test data shape: [10000, 28, 28]

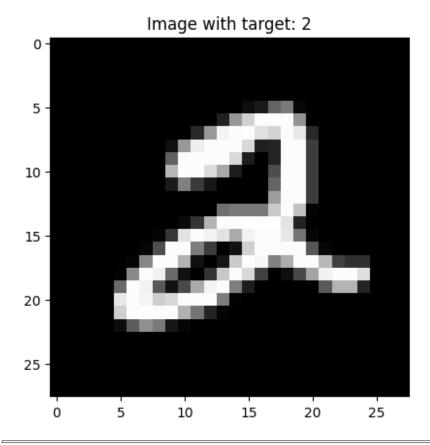
```
# Show an example image
plt.imshow(train_data.data[5], cmap="gray")
plt.title("Image with target: %i" % train_data.targets[5])

Text(0.5, 1.0, 'Image with target: 2')

Text(0.5, 1.0, 'Image with target: 2')
```







Creating a CNN

Convolutional neural networks (CNNs) are popular with image data. The network automatically extracts useful features from images, such as edges, contours, and objects.

This lab introduces CNNs, but the details of CNNs will be discussed in a later module.

CNNs require minimal preprocessing compared to older algorithms, such as feed-forward neural networks, that are used for computer vision. Although feed-forward neural networks can still be used with image data, CNNs can capture the spatial and temporal properties in an image with a significant reduction in the number of parameters. In this notebook, you will use a simple CNN to extract information from image data.

You will use PyTorch's Conv2D layer with the following interface to process the images:

nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, ...)

Parameter definitions:

- in_channels (int): Number of channels in the input image
- out_channels (int): Number of channels that are produced by the convolution
- **kernel_size** (int or tuple): Size of the convolving kernel
- **stride (int or tuple, optional):** Stride of the convolution (default is 1)

The output dimension of the Conv2D layer can be calculated using the following formula:

```
((W - K + 2P)/S + 1)
```

Where:

- W = Input size
- K = Kernel size
- S = Stride
- P = Padding (not used in the notebook)

Example:

```
For an image of size = (28x28), kernel size = 3, stride = 1, and padding = 0, the output dimension is (28 - 3 + 0)/1 + 1 = 26.
```

With out channels = 1, the output dimension is (26, 26).

With out channels = 3, the output dimension is (26, 26, 3).

```
# Define hyperparameters
batch_size = 100 # Size of input data for one iteration
num classes = 10 # Number of output classes, discrete range [0,9]
num epochs = (
   10 # Number of times that the entire dataset is passed through
the model
# Size of step
lr = 1e-3
# Use GPU if available; otherwise, use CPU
device = torch.device("cuda:0" if torch.cuda.is available() else
"cpu")
# Use PyTorch DataLoaders to load the data in batches
train loader = torch.utils.data.DataLoader(
   dataset=train data, batch size=batch size, shuffle=True,
drop last=True
# Repeat for test dataset
test loader = torch.utils.data.DataLoader(
   dataset=test data, batch size=batch size, shuffle=False
input size = 26 * 26 * 32 # Flattened dimension for the linear layer
net = nn.Sequential(
   nn.Conv2d(in channels=1, out channels=32, kernel size=3),
```

```
nn.ReLU(),
   nn.Flatten(),
   nn.Linear(26 * 26 * 32, 128),
   nn.ReLU(),
   nn.Linear(128, num classes),
   nn.Softmax(dim=1)
)
net = net.to(device)
def xavier init weights(m):
   if type(m) == nn.Linear:
       torch.nn.init.xavier uniform (m.weight)
# Initialize weights/parameters for the network
net.apply(xavier init weights)
Sequential(
  (0): Conv2d(1, 32, kernel size=(3, 3), stride=(1, 1))
  (1): ReLU()
  (2): Flatten(start dim=1, end dim=-1)
  (3): Linear(in features=21632, out features=128, bias=True)
  (4): ReLU()
  (5): Linear(in features=128, out features=10, bias=True)
  (6): Softmax(dim=1)
)
Sequential(
  (0): Conv2d(1, 32, kernel size=(3, 3), stride=(1, 1))
  (1): ReLU()
  (2): Flatten(start dim=1, end dim=-1)
  (3): Linear(in features=21632, out features=128, bias=True)
  (4): ReLU()
  (5): Linear(in features=128, out features=10, bias=True)
  (6): Softmax(dim=1)
)
# Define the loss function and the optimizer
# Choose cross-entropy loss for this classification problem
loss = nn.CrossEntropyLoss()
# Choose the Adam optimizer. You can also experiment with other
optimizers.
optimizer = torch.optim.Adam(net.parameters(), lr=lr)
```

Training the network

Now you are ready to train the CNN.

```
import time
# Network training and validation
# Start the outer epoch loop (epoch = full pass through the dataset)
for epoch in range(num epochs):
    start = time.time()
    training loss = 0.0
    # Training loop (with autograd and trainer steps)
    # This loop trains the neural network
    # Weights are updated here
    net.train() # Activate training mode (dropouts and so on)
    for images, target in train_loader:
        # Zero the parameter gradients
        optimizer.zero grad()
        images = images.to(device)
        target = target.to(device)
        # Forward + backward + optimize
        output = net(images)
        L = loss(output, target)
        L.backward()
        optimizer.step()
        # Add batch loss
        training loss += L.item()
    # Take the average losses
    training loss = training loss / len(train loader)
    end = time.time()
    print("Epoch %s. Train loss %f Seconds %f" % (epoch,
training loss, end - start))
Epoch O. Train loss 2.018384 Seconds 14.523459
Epoch 0. Train_loss 2.018384 Seconds 14.523459
Epoch 1. Train loss 1.556107 Seconds 6.641848
Epoch 1. Train loss 1.556107 Seconds 6.641848
Epoch 2. Train loss 1.484114 Seconds 6.657332
Epoch 2. Train loss 1.484114 Seconds 6.657332
Epoch 3. Train loss 1.477806 Seconds 6.642189
Epoch 3. Train loss 1.477806 Seconds 6.642189
Epoch 4. Train_loss 1.473900 Seconds 6.661234
Epoch 4. Train loss 1.473900 Seconds 6.661234
Epoch 5. Train loss 1.471776 Seconds 6.611521
Epoch 5. Train loss 1.471776 Seconds 6.611521
```

```
Epoch 6. Train_loss 1.469788 Seconds 6.640481
Epoch 6. Train_loss 1.469788 Seconds 6.640481
Epoch 7. Train_loss 1.468814 Seconds 6.646728
Epoch 7. Train_loss 1.468814 Seconds 6.646728
Epoch 8. Train_loss 1.467432 Seconds 6.641755
Epoch 8. Train_loss 1.467432 Seconds 6.641755
Epoch 9. Train_loss 1.466805 Seconds 6.643712
Epoch 9. Train_loss 1.466805 Seconds 6.643712
```

Testing the network

Finally, evaluate the performance of the trained network on the test set.

```
from sklearn.metrics import classification report
net.eval() # Activate eval mode (don't use dropouts and such)
# Get test predictions
predictions, labels = [], []
for images, target in test loader:
    images = images.to(device)
    target = target.to(device)
    predictions.extend(net(images).argmax(axis=1).tolist())
    labels.extend(target.tolist())
# Print performance on the test data
print(classification report(labels, predictions, zero division=1))
                            recall f1-score
              precision
                                                support
           0
                   0.99
                              0.99
                                        0.99
                                                    980
           1
                   0.99
                              0.99
                                        0.99
                                                   1135
           2
                   0.98
                              0.98
                                        0.98
                                                   1032
           3
                   0.98
                              0.99
                                        0.99
                                                   1010
           4
                              0.99
                   0.98
                                        0.99
                                                    982
           5
                   0.98
                              0.98
                                        0.98
                                                    892
           6
                   0.99
                              0.99
                                        0.99
                                                    958
           7
                   0.96
                              0.99
                                        0.97
                                                   1028
           8
                   0.99
                              0.96
                                        0.98
                                                    974
           9
                   0.99
                              0.97
                                        0.98
                                                   1009
    accuracy
                                        0.98
                                                  10000
                   0.98
                              0.98
                                        0.98
                                                  10000
   macro avg
                                        0.98
weighted avg
                   0.98
                              0.98
                                                  10000
              precision
                            recall f1-score
                                                support
```

0	0.99	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.98	0.98	0.98	1032
3	0.98	0.99	0.99	1010
4	0.98	0.99	0.99	982
5	0.98	0.98	0.98	892
6	0.99	0.99	0.99	958
7	0.96	0.99	0.97	1028
8	0.99	0.96	0.98	974
9	0.99	0.97	0.98	1009
accuracy			0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000

Conclusion

In this notebook, you practiced using a CNN.

Next Lab: Processing text

In the next lab you will learn how to do more advanced text processing.