

```
#!/usr/bin/env python
# coding: utf-8

# <center></center> <br/>
#
# # Application of Deep Learning to Text and Image Data
# ## Module 1, Lab 1, Notebook 2: Examining a Neural Network Architecture
#
# # In this notebook, you will implement a minimum viable neural network to see the different architecture components.
#
# # The simplest possible neural network architecture is logistic regression. This lab will cover data ingestion, how to define the model, loss function, and the
#
# # You will do the following:
#
# # - Generate a simulated dataset
# # - Use basic components of a neural network
# # - Implement a neural network by using PyTorch
# # - Train a neural network
#
# ---
#
# # You will be presented with activities throughout the notebook: <br/>
#
# ||
# | --- |
# |<p style="text-align:center;"> No coding is needed for an activity. You try to understand a concept, <br/>answer questions, or run a code cell.</p>|
#
# ---
#
# ## Index
#
# * [Simulated dataset] (#Simulated-dataset)
# * [Neural network basics] (#Neural-network-basics)
# * [Implementing a neural network with PyTorch] (#Implementing-a-neural-network-with-PyTorch)
# * [Training of the neural network] (#Training-of-the-neural-network)
#
# ---
#
# ## Simulated dataset
#
# # In this example, you will train a neural network on a dataset that is randomly generated. The dataset will have two classes, and you will train the neural ne
#
# In[1]:
#
# Install libraries
get_ipython().system('pip install -U -q -r requirements.txt')
#
# In[2]:
#
# Load the sample data
from sklearn.datasets import make_circles
#
# Specify settings, including how many examples to extract
X, y = make_circles(
    n_samples=750, shuffle=True, random_state=42, noise=0.05, factor=0.3
)
#
# Plot the simulated dataset.
#
# In[ ]:
#
get_ipython().run_line_magic('matplotlib', 'inline')
import matplotlib.pyplot as plt
import seaborn as sns

def plot_dataset(X, y, title):
    # Activate the Seaborn visualization
    sns.set()

    # Plot both classes: Class 1 is blue, Class 2 is red
    plt.scatter(X[y == 1, 0], X[y == 1, 1], c="blue", label="Class 1")
    plt.scatter(X[y == 0, 0], X[y == 0, 1], c="red", label="Class 2")
    plt.legend(loc="upper right")
    plt.xlabel("x1")
    plt.ylabel("x2")
    plt.xlim(-2, 2)
    plt.ylim(-2, 2)
    plt.title(title)
    plt.show()

plot_dataset(X, y, title="Dataset")

# The goal is to build a neural network that can differentiate between the two classes in the dataset. The simplest neural network that can tackle this problem
#
# In[ ]:
#
import torch

# Use a GPU resource if available; otherwise, use CPU
device = "cuda" if torch.cuda.is_available() else "cpu"

# To work with PyTorch, you need to convert the dataset to tensors first
X = torch.tensor(X, dtype=torch.float32).to(device)
y = torch.tensor(y, dtype=torch.float32).reshape((-1, 1)).to(device)

# ---
```

```

# ## Neural network basics
#
# The fundamental building blocks of neural networks are *neurons*, which are functions that are followed by an activation function. It's common for the initial
# 
#
# Your first goal is to implement a simple neural network with one neuron that uses the sigmoid function as the activation function to predict class 1 or class
# As an equation, this would look like the following: For some input  $\mathbf{X}$  and output  $\mathbf{y}$ , the logistic regression model is defined as:
#
# 
$$\hat{y} = \sigma(\mathbf{X}\mathbf{w} + \mathbf{b})$$

#
# with some initial choices for the parameters,  $\mathbf{w}$  weights matrix and bias  $\mathbf{b}$ .
#
# You don't know what the best values for  $\mathbf{w}$  and  $\mathbf{b}$  are, so initialize the weights matrix  $\mathbf{w}$  at random with zero mean and standard deviation
#
# In [ ]:

# Define the logistic regression
def log_reg(X, w, b):
    return torch.sigmoid(torch.matmul(X, w) + b)

# Define the basic neural network
net = log_reg

# Initialize for w and b
w = torch.normal(0, 1, size=(2, 1), requires_grad=True).to(device)
b = torch.zeros(1, requires_grad=True).to(device)

# To test the neural net, pass in an example data point
print(
    f"For datapoint 0, the probability of being class 1 is {float(net(X[0], w, b).item()):.2f}."
)

# This is a basic single-layer neural network that hasn't been trained yet. Ultimately, the goal is to find the best possible values for  $\mathbf{w}$  and  $\mathbf{b}$ 
#
# In [ ]:

# Import system library and append path
import sys

sys.path.insert(1, "..")

# Import utility functions that provide answers to challenges
from MLUDTI_EN_M1_Lab1_quiz_questions import *

# <div style="border: 4px solid coral; text-align: center; margin: auto;">
# <h3><i>Try it yourself!</i></h3>
# <br>
# <p style="text-align:center;margin:auto;"> </p>
# <p style="text-align: center; margin: auto;">It's time to check your knowledge. To load the question, run the following cell.</p>
# <br>
# </div>

# In [ ]:

# Run this cell to display the question and check your answer
question_4

#
# ---
# ## Implementing a neural network with PyTorch
#
# Now you need to build, train, and validate a neural network in PyTorch. With PyTorch, you can list the different layers and activation functions that you want
#
# You can use the `Sequential()` function to define the functions in the order that you want them to run. The first row refers to the first function to be used
#
# In [ ]:

from torch import nn

# Create a sequential container that chains a linear regression function with a sigmoid activation function
net = nn.Sequential(
    nn.Linear(2, 1), # Linear layer-1 with 1 out_features and input size 2
    nn.Sigmoid(), # Sigmoid activation function
)

# ### Initialization
#
# Before you continue, you need to initialize PyTorch values. This is important for the same reason that you initialized values for  $\mathbf{w}$  and  $\mathbf{b}$ 
#
# Picking the starting point is critical. Researchers have developed several initialization strategies that you can use. The _Xavier initialization_ is common
#
# For a list of initializers, see the [PyTorch documentation](https://pytorch.org/docs/stable/nn.init.html). For information about Xavier initialization, see [
#
# In [ ]:

# Initialize the weights with an Xavier initializer
def xavier_init_weights(m):
    if type(m) == nn.Linear:
        torch.nn.init.xavier_uniform_(m.weight)
        torch.nn.init.zeros_(m.bias)

```

```

# Apply the initialization to the sequential network that you created earlier
net.apply(xavier_init_weights)

### Loss function
#
# Now that you have set up a neural network, you need to select a loss function to quantify how good a given selection of parameters are. Many loss functions e
#
# Binary cross-entropy loss (log loss) is a loss function that is commonly used for binary classification:
#
# ```python
# loss = nn.BCELoss()
# ```
#
# During the training of the neural network, the initial model parameters will be updated until model predictions fit the data sufficiently well. One way to co
#
# For a full list of supported loss functions, see [torch.nn Loss Functions](https://pytorch.org/docs/stable/nn.html#loss-functions) in the PyTorch documentati

### Optimization method
# Each update requires taking the gradient of the loss function with respect to the parameters. Automatic differentiation is used to compute the gradient, and
#
# The `torch.optim` module provides necessary optimization algorithms for neural networks. You can use an optimizer to train a network by using the stochastic
#
# ```python
# from torch import optim
# optimizer = optim.SGD(net.parameters(), lr=0.001)
# ```
#
# Three lines of code are required to perform a gradient descent update:
#
# ```
# loss.backward() # Compute updates for each parameter
# optimizer.step() # Make the updates for each parameter
# optimizer.zero_grad() # Clean-up step for PyTorch
# ```

# ---
### Training of the neural network
#
# Now that you have all the components, you can create a loop that takes a given set of parameter values, creates outputs, evaluates the performance, and updat

# Start by looking at the basic network again.

In [ ]:

# Print network
net.to(device)

# Next, specify the loss function, optimization method, and how many epochs (loops) are run to update the parameters.

In [ ]:

num_epochs = 50 # Total number of epochs (loops)

# Define the loss. Because you used sigmoid in the last layer, use nn.BCELoss.
# Otherwise, you could have used nn.BCEWithLogitsLoss. This loss combines a sigmoid layer and the BCELoss in a single class.
loss = nn.BCELoss(reduction="none")

# Define the optimizer, SGD with learning rate
optimizer = torch.optim.SGD(net.parameters(), lr=0.01)

# Finally, it is time for training!
#
# Training will run through the dataset 50 times, and print training and validation losses after each epoch.

In [ ]:

train_losses = []
for epoch in range(num_epochs):
    training_loss = 0
    # Zero the parameter gradients
    optimizer.zero_grad()
    output = net(X)
    L = loss(output, y).sum()
    training_loss += L.item()
    L.backward()
    optimizer.step()
    training_loss = training_loss / len(y)
    train_losses.append(training_loss)

# Now that training has completed, you can plot the training and validation loss plots.

In [ ]:

get_ipython().run_line_magic('matplotlib', 'inline')
import matplotlib.pyplot as plt
import seaborn as sns

plt.plot(train_losses, label="Training Loss")
plt.title("Loss vs. Epoch")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()

# Notice that the loss (errors) decreases as the training process continues, as expected.
#
# One final step is to compare this plot to the validation loss to see if it is overfitting the model.

```

```
# ---
# ## Conclusion
#
# In this notebook, you learned about training a neural network. You should now understand the basic steps of building a neural network and how to evaluate its
#
# ---
# ## Next lab
# In the next lab, you will learn about the multilayer perceptron, which is the simplest feed-forward neural network architecture, and how to use dropout layer
```